

# MATH 173A: Homework #2

Due on Oct 22, 2024 at 23:59pm

*Professor Cloninger*

**Ray Tsai**

A16848188

## Problem 1

Using the conditions of optimality, find the extreme points of the following functions and determine whether they are maxima or minima. You may use a computer to find the eigenvalues, but these questions should have easily accessible eigenvalues by hand.

(a)  $f : \mathbb{R}^2 \rightarrow \mathbb{R}$  for  $f(x_1, x_2) = x_1^4 + 2x_2^4 - 4x_1x_2$

*Proof.* Note that

$$\begin{aligned}\nabla f(x) &= (4x_1^3 - 4x_2, 8x_2^3 - 4x_1) = 0 \\ \nabla^2 f(x) &= \begin{bmatrix} 12x_1^2 & -4 \\ -4 & 24x_2^2 \end{bmatrix}.\end{aligned}$$

Thus, the critical points are  $x^* = (0, 0)$  or  $\pm(2^{-1/8}, 2^{-3/8})$ . We can then check

$$\begin{aligned}\nabla^2 f(0, 0) &= \begin{bmatrix} 0 & -4 \\ -4 & 0 \end{bmatrix} \\ \nabla^2 f(2^{-1/8}, 2^{-3/8}) &= \nabla^2 f(-2^{-1/8}, -2^{-3/8}) = \begin{bmatrix} 12 \cdot 2^{-1/4} & -4 \\ -4 & 24 \cdot 2^{-3/4} \end{bmatrix}.\end{aligned}$$

Since the eigenvalues of  $\nabla^2 f(0, 0)$  are  $\pm 4$ ,  $(0, 0)$  is a saddle point. Since  $\det \nabla^2 f(2^{-1/8}, 2^{-3/8}) = \det \nabla^2 f(-2^{-1/8}, -2^{-3/8}) = 128 > 0$  and  $\frac{\partial^2 f}{\partial x_1^2} > 0$ , the critical points  $\pm(2^{-1/8}, 2^{-3/8})$  are local minima.  $\square$

(b)  $f : \mathbb{R}^3 \rightarrow \mathbb{R}$  for  $f(\vec{x}) = \vec{x}^T A \vec{x} + b^T \vec{x}$ , where

$$A = \begin{bmatrix} -1 & 0 & \frac{1}{2} \\ 0 & -1 & 0 \\ \frac{1}{2} & 0 & -1 \end{bmatrix}, \quad b = \begin{bmatrix} 0 \\ 1 \\ 2 \end{bmatrix}$$

*Proof.*

$$\begin{aligned}\nabla f(\vec{x}) &= 2A\vec{x} + b, \\ \nabla^2 f(\vec{x}) &= 2A.\end{aligned}$$

Setting  $\nabla f(\vec{x}) = 0$  yields

$$\vec{x}^* = -\frac{1}{2}A^{-1}b = \begin{bmatrix} \frac{2}{3} \\ \frac{1}{2} \\ \frac{4}{3} \end{bmatrix}.$$

Since the eigenvalues of  $\nabla^2 f(\vec{x}) = 2A$  are  $-1, -2, -3$ , we know  $\nabla^2 f(\vec{x}) \prec 0$  and the critical point is a local maximum.  $\square$

## Problem 2

Consider the problem  $f : \mathbb{R}^n \rightarrow \mathbb{R}$  for  $f(x) = \|Ax - b\|_2^2$  for  $A \in \mathbb{R}^{m \times n}$  and  $b \in \mathbb{R}^m$ . (Note: this was on the last homework). Write down the gradient descent algorithm to solve the optimization

$$\min_{x \in \mathbb{R}^n} f(x).$$

*Proof.* We already know  $f$  is convex and the gradient of  $f$  is

$$\nabla f(x) = 2A^T(Ax - b).$$

Thus, the gradient descent algorithm is

$$\begin{aligned} x^{(0)} &= \text{any vector from } \mathbb{R}^n \\ x^{(t+1)} &= x^{(t)} - \mu^{(t)} \nabla f(x^{(t)}) = x^{(t)} - 2\mu^{(t)} A^T(Ax^{(t)} - b) \end{aligned}$$

where  $\mu^{(t)} > 0$  and the terminating condition is of our choice. □

## Problem 3

**Implementing Classification Model:** First some background for classification:

- You are given labeled data  $\{(x_i, y_i)\}_{i=1}^N$  for  $x_i \in \mathbb{R}^d$  and  $y_i \in \{-1, 1\}$ .
- Logistic regression involves choosing a label according to

$$y = \text{sign}(\langle w, x \rangle).$$

Note we ignore the  $y$ -intercept term here, so we only need the optimal  $w \in \mathbb{R}^d$ .

- It turns out the correct function to minimize to find the weights is

$$F(w) = \frac{1}{N} \sum_{i=1}^N \log(1 + e^{-\langle w, x_i \rangle y_i}).$$

- (a) Is  $F(w)$  a convex function?

*Proof.* By the chain rule,

$$\begin{aligned} \nabla F(w) &= \frac{1}{N} \sum_{i=1}^N \frac{-y_i e^{-\langle w, x_i \rangle y_i}}{1 + e^{-\langle w, x_i \rangle y_i}} x_i \\ \nabla^2 F(w) &= \frac{1}{N} \sum_{i=1}^N \frac{y_i^2 e^{-\langle w, x_i \rangle y_i}}{(1 + e^{-\langle w, x_i \rangle y_i})^2} x_i x_i^T. \end{aligned}$$

Since  $x_i x_i^T \succeq 0$ ,  $\frac{y_i^2 e^{-\langle w, x_i \rangle y_i}}{(1 + e^{-\langle w, x_i \rangle y_i})^2} \geq 0$ , and the sum of positive semidefinite matrices is positive semidefinite,  $\nabla^2 F(w) \succeq 0$ . Thus,  $F(w)$  is convex.  $\square$

- (b) Find a gradient descent algorithm for minimizing  $F$ .

*Proof.* The gradient descent algorithm is

$$\begin{aligned} w^{(0)} &= \text{any vector from } \mathbb{R}^d \\ w^{(t+1)} &= w^{(t)} - \mu^{(t)} \nabla F(w^{(t)}) = w^{(t)} + \frac{\mu^{(t)}}{N} \sum_{i=1}^N \frac{y_i e^{-\langle w^{(t)}, x_i \rangle y_i}}{1 + e^{-\langle w^{(t)}, x_i \rangle y_i}} x_i \end{aligned}$$

where  $\mu^{(t)} > 0$  and the terminating condition is of our choice.  $\square$

## Problem 4

**Coding Question:** Recall that the equation for an ellipse in  $\mathbb{R}^2$  is

$$a_1x^2 + a_2y^2 = 1.$$

Given data  $\{(x_i, y_i)\}_{i=1}^N \subset \mathbb{R}^2$  that lie on (or near) the ellipse, you can find the best fit ellipse by solving

$$\min_{\mathbf{a} \in \mathbb{R}^2} f(\mathbf{a})$$

where

$$f(\mathbf{a}) = \sum_{i=1}^N (a_1x_i^2 + a_2y_i^2 - 1)^2.$$

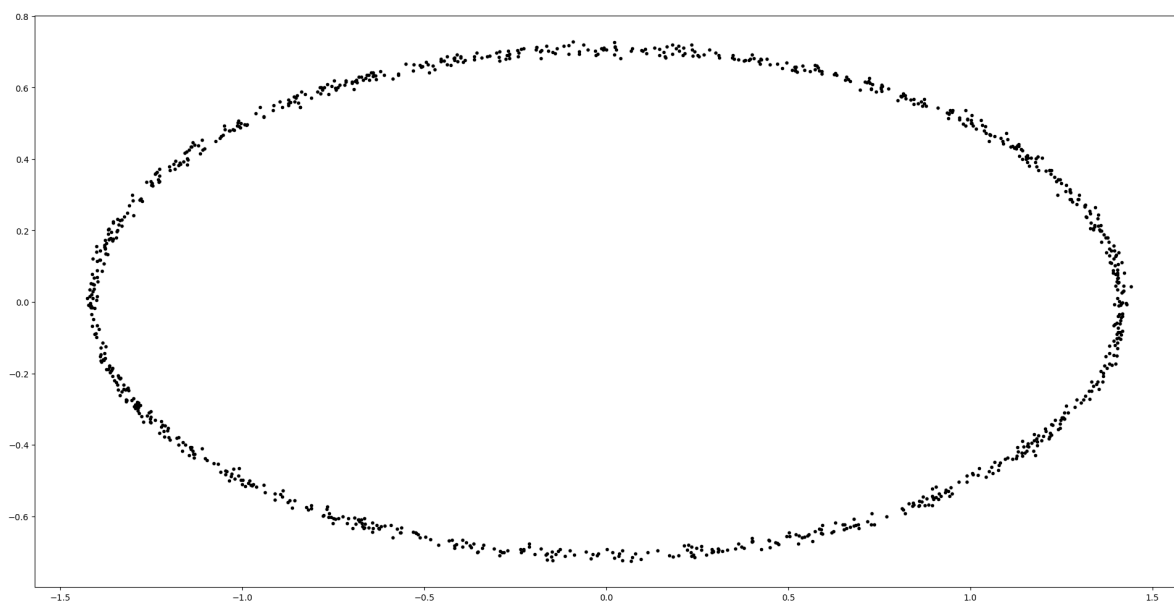
(a) Find an  $A \in \mathbb{R}^{N \times 2}$  and  $b \in \mathbb{R}^N$  such that  $f(\mathbf{a}) = \|A\mathbf{a} - b\|_2^2$ . What is  $A$  in terms of  $(x_i, y_i)$ ?

*Proof.* Put  $A = \begin{bmatrix} x_1^2 & y_1^2 \\ \vdots & \vdots \\ x_N^2 & y_N^2 \end{bmatrix}$  and  $b = \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix}$  and we have

$$\begin{aligned} f(\mathbf{a}) &= \sum_{i=1}^N (a_1x_i^2 + a_2y_i^2 - 1)^2 \\ &= \sum_{i=1}^N \left( \begin{bmatrix} x_i^2 & y_i^2 \end{bmatrix} \begin{bmatrix} a_1 \\ a_2 \end{bmatrix} - 1 \right)^2 \\ &= \|A\mathbf{a} - b\|_2^2. \end{aligned}$$

□

(b) Download the data provided on the HW page (called `HW2.ellipse.csv`), and create a scatter plot of the points (submit your code and the plot).



**Code:**

```
ellipse = pd.read_csv('HW2-ellipse.csv', header=None) ellipse.columns = ['x', 'y']
ax.plot(ellipse['x'], ellipse['y'], '.', color='black')
```

- (c) Using Problem 3, create computer code to compute the gradient descent algorithm on this  $f(\mathbf{a})$ . The code must include a stopping condition. Use a step-size of  $\mu = \frac{1}{2\|\mathbf{A}^T\mathbf{A}\|}$ . Note, you cannot use a built-in gradient descent algorithm; it must be written with a while or for loop. Also note, the norm of a matrix  $\|\mathbf{X}\| = \lambda_{\max}(\mathbf{X})$  is the largest eigenvalue of  $\mathbf{X}$ , and can be computed using `norm(X,2)` in MATLAB or `np.linalg.norm(X,2)` in Python. (Submit the code)

**Code:**

```
A = ellipse.values * ellipse.values b = np.ones((A.shape[0], 1)) mu = 1/(2 * np.linalg.norm(A.T @ A))

def df(x): return 2 * A.T @ (A @ x - b)

for i in range(1000): a = a - mu * df(a)
```

- (d) Using the data provided and your gradient descent code, estimate the solution  $\mathbf{a}$ . Report  $\mathbf{a}$  and  $f(\mathbf{a})$ . Given  $f(\mathbf{a})$  and  $N$ , do you think you fit the data well or poorly? Given the convexity of  $f$ , do you think this is the optimal  $\mathbf{a}$ ?

*Solution.* The solution estimated by the code is  $\mathbf{a} \approx \begin{bmatrix} 0.5001 \\ 1.9946 \end{bmatrix}$ , with  $f(\mathbf{a}) \approx 0.4641$ . Since we are given  $N = 1000$  points, on average each point is only off by  $f(\mathbf{a})/N \approx 0.0004641$ , which is a small enough number to me. Given that  $\Delta f(\mathbf{a}) \approx 0$ , I believe  $\mathbf{a}$  has reached a local minimum up to some computational error. Since  $f$  is convex,  $\mathbf{a}$  is extremely close to the theoretical optimal solution.  $\square$