

Question 4

Let $A \in \mathbb{R}^{n \times n}$ be a diagonal matrix with diagonal entries

$$A_{ii} = i, \quad \text{i.e. the entries run from 1 to } n,$$

and let $b \in \mathbb{R}^n$ a vector with all 1 entries. Define the function

$$f(x) = \frac{1}{2}x^T A x - b^T x.$$

We want to compare the convergence behavior of conjugate gradient (version 0 or 1) and gradient descent. Do the following for $n = 20$ and $n = 100$ with initialization $x^{(0)} = 0$.

In [122... `import numpy as np`
`from matplotlib import pyplot as plt`

Part A

Find the optimal solution x^* by solving $Ax = b$ using a Matlab/Python linear equation solver (or by hand and hard code the answer).

In [123... `n = 20`

```
def A(n):
    return np.diag(np.arange(1, n+1))

def b(n):
    return np.ones(n)

def x_opt(n):
    return np.linalg.solve(A(n), b(n))
```

Part B

Program and run the gradient descent method for f with a fixed stepsize. Run the method for n iterations. You may experiment with the stepsize until you see something that works or use a stepsize dictated by a theorem in the class.

```
In [124...
def f(x, n):
    return 1/2 * x.T @ A(n) @ x - b(n) @ x

def df(x, n):
    return A(n) @ x - b(n)

def gd(x, n, mu = 2e-2):
    return x - mu * df(x, n)
```

```
In [125... N = [20, 100]

gd_x_values = [], []
gd_f_values = [], []

for n in N:
    x = np.zeros(n)
    for i in range(n):
        gd_x_values[N.index(n)].append(np.linalg.norm(x - x_opt(n)))
        gd_f_values[N.index(n)].append(f(x, n) - f(x_opt(n), n))
    x = gd(x, n)
```

Part C

Program and run the conjugate gradient (version 0 or 1) for f . Run the method for n iterations.

```
In [126... N = [20, 100]

cgd_x_values = [], []
cgd_f_values = [], []

for n in N:
    x = np.zeros(n)
    r = df(x, n)
    p = -r

    for i in range(n):
        cgd_x_values[N.index(n)].append(max(np.linalg.norm(x - x_opt(n)), 1e-16))
        cgd_f_values[N.index(n)].append(max(f(x, n) - f(x_opt(n), n), 1e-16))

        alpha = r.T @ r / (p.T @ A(n) @ p)
        x += alpha * p
        r_new = r + alpha * A(n) @ p
        beta = r_new.T @ r_new / (r.T @ r)
        p = -r_new + beta * p
        r = r_new
```

Plot the $f(x^{(t)}) - f(x^*)$ for both methods in the same figure. In a different figure, plot $\|x^{(t)} - x^*\|$ for both methods. If you encounter a number smaller than 10^{-16} , set it to be 10^{-16} . In both plots, make the logarithmic scale for the vertical axis. Comment on the plots.

```
In [127... plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(range(20), gd_f_values[0])
plt.yscale('log')
plt.xlabel(f"Iterations")
plt.ylabel(r"Value of $log [f(x) - f(x^*)]$" )
plt.title(f"n = 20")

plt.subplot(1, 2, 2)
plt.plot(range(20), cgd_f_values[0])
plt.yscale('log')
plt.xlabel(f"Iterations")
plt.ylabel(r"Value of $log [f(x) - f(x^*)]$" )
plt.title(f"n = 20")

plt.show()

plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(range(20), gd_x_values[0])
plt.yscale('log')
plt.xlabel(f"Iterations")
plt.ylabel(r"Value of $log \|x - x^*\| $" )
plt.title(f"n = 20")

plt.subplot(1, 2, 2)
plt.plot(range(20), cgd_x_values[0])
plt.yscale('log')
plt.xlabel(f"Iterations")
plt.ylabel(r"Value of $log \|x - x^*\| $" )
plt.title(f"n = 20")

plt.show()

plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(range(100), gd_f_values[1])
plt.yscale('log')
plt.xlabel(f"Iterations")
plt.ylabel(r"Value of $log [f(x) - f(x^*)]$" )
```

```

plt.title(f"n = 100")

plt.subplot(1, 2, 2)
plt.plot(range(100), cgd_f_values[1])
plt.yscale('log')
plt.xlabel(f"Iterations")
plt.ylabel(r"Value of $log [f(x) - f(x^*)]$")
plt.title(f"n = 100")

plt.show()


plt.figure(figsize=(12, 5))

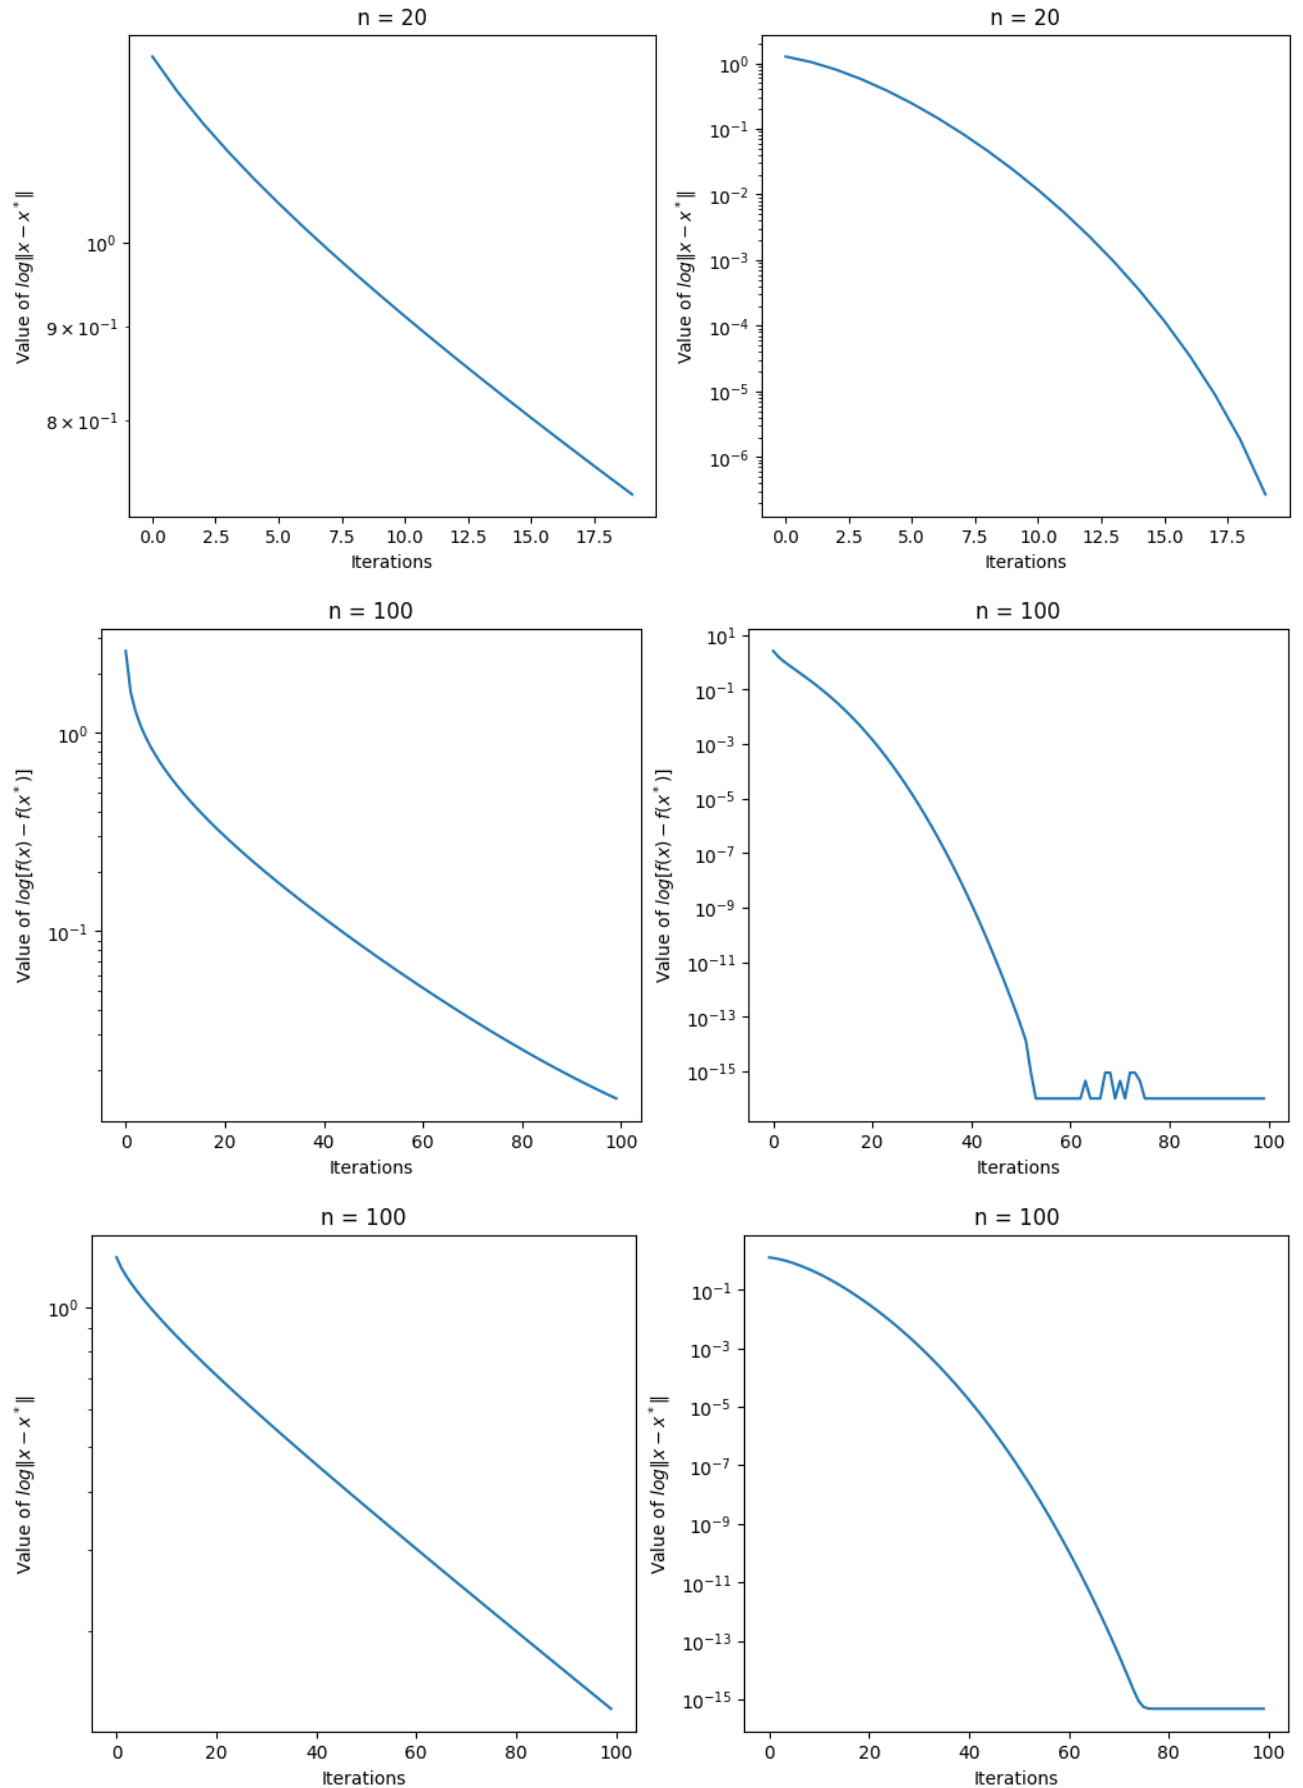
plt.subplot(1, 2, 1)
plt.plot(range(100), gd_x_values[1])
plt.yscale('log')
plt.xlabel(f"Iterations")
plt.ylabel(r"Value of $log \||x - x^*|$")
plt.title(f"n = 100")

plt.subplot(1, 2, 2)
plt.plot(range(100), cgd_x_values[1])
plt.yscale('log')
plt.xlabel(f"Iterations")
plt.ylabel(r"Value of $log \||x - x^*|$")
plt.title(f"n = 100")

plt.show()

```





The conjugate gradient descent method converges significantly faster than the standard gradient descent. The conjugate gradient descent method indeed converges within n iterations, agreeing with the theorem we learned.