

Question 4

We will implement the SVM algorithm with gradient descent to classify two gaussians in 2D. The dataset is given in `HW7Q4.csv`.


```
In [33]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

Part A

In `HW7Q4.csv`, the first 100 rows are the data for cluster 1: $(x_i, y_i) \in \mathbb{R}^2 \times \mathbb{R}$, $i = 1, \dots, 100$, with $y_i = 1$ always. The next 100 rows are the data for cluster 2: $(x_i, y_i) \in \mathbb{R}^2 \times \mathbb{R}$, $i = 101, \dots, 200$, with $y_i = -1$ always. Create and turn in a scatter plot of the feature vectors, i.e., the x_i 's, colored by the label, i.e., y_i 's (blue for 1 and red for -1).

```
In [34]: data = pd.read_csv('HW7Q4.csv', header=None)
data.columns = ['x1', 'x2', 'y']

plt.figure(figsize=(8, 6))
plt.scatter(data[:100]['x1'], data[:100]['x2'], color='blue', label='y=1')
plt.scatter(data[100:]['x1'], data[100:]['x2'], color='red', label='y=-1')
plt.legend()
plt.show()
```



Part B

Create a function for the gradient of the loss

$$L(w) = \frac{1}{2} \|w\|^2 + \sum_{i=1}^n \max(0, 1 - y_i \langle x_i, w \rangle)$$

$$\nabla L(w) = w + \sum_{i=1}^n -y_i x_i \cdot \mathbf{1}_{1 - y_i \langle x_i, w \rangle > 0},$$

where $\mathbf{1}_{1 - y_i \langle x_i, w \rangle > 0} = \begin{cases} 1 & \text{if } 1 - y_i \langle x_i, w \rangle > 0 \\ 0 & \text{else} \end{cases}$. Also, here $n = 200$. To compute the gradient, you'll have to compute an indicator of whether $1 - y_i \langle x_i, w \rangle$ is positive or negative at every point, and sum up the contribution of this term for all points where it's positive.

```
In [43]: n = len(data)
```

```
def indicator(w, x1, x2, y):
    if 1 - y * (w[0] * x1 + w[1] * x2) > 0:
        return 1
    return 0

def L(w):
    return 1/2 * np.linalg.norm(w)**2 + sum([max(0, 1 - row['y'] * (w[0] * r

def dL(w):
    return w + sum([-row['y'] * np.array([row['x1'], row['x2']]) * indicator
```

Part C

Setting the step size $\mu = 10^{-4}$ and starting at $w^{(0)} = (-1, 1)$, run 1000 iterations of gradient descent. You will create two plots.

- Plot the classification error (averaged over all the points) as a function of the iterations. The classification of x_i is determined by $\text{sign}(\langle x_i, w \rangle)$.
- Plot the margin $\frac{2}{\|w\|}$ as a function of the iterations. This shows how much of a gap you have between the classes you've learned.

```
In [ ]: T = 1000
mu = 1e-4

def error(w):
    sum = 0
    for _, row in data.iterrows():
        if np.sign(w[0] * row['x1'] + w[1] * row['x2']) != row['y']:
            sum += 1
    return sum / n

w = np.array([-1, 1])
error_val = []
margin_val = []


for t in range(T):
    w = w - mu * dL(w)
    error_val.append(error(w))
    margin_val.append(2 / np.linalg.norm(w))
```

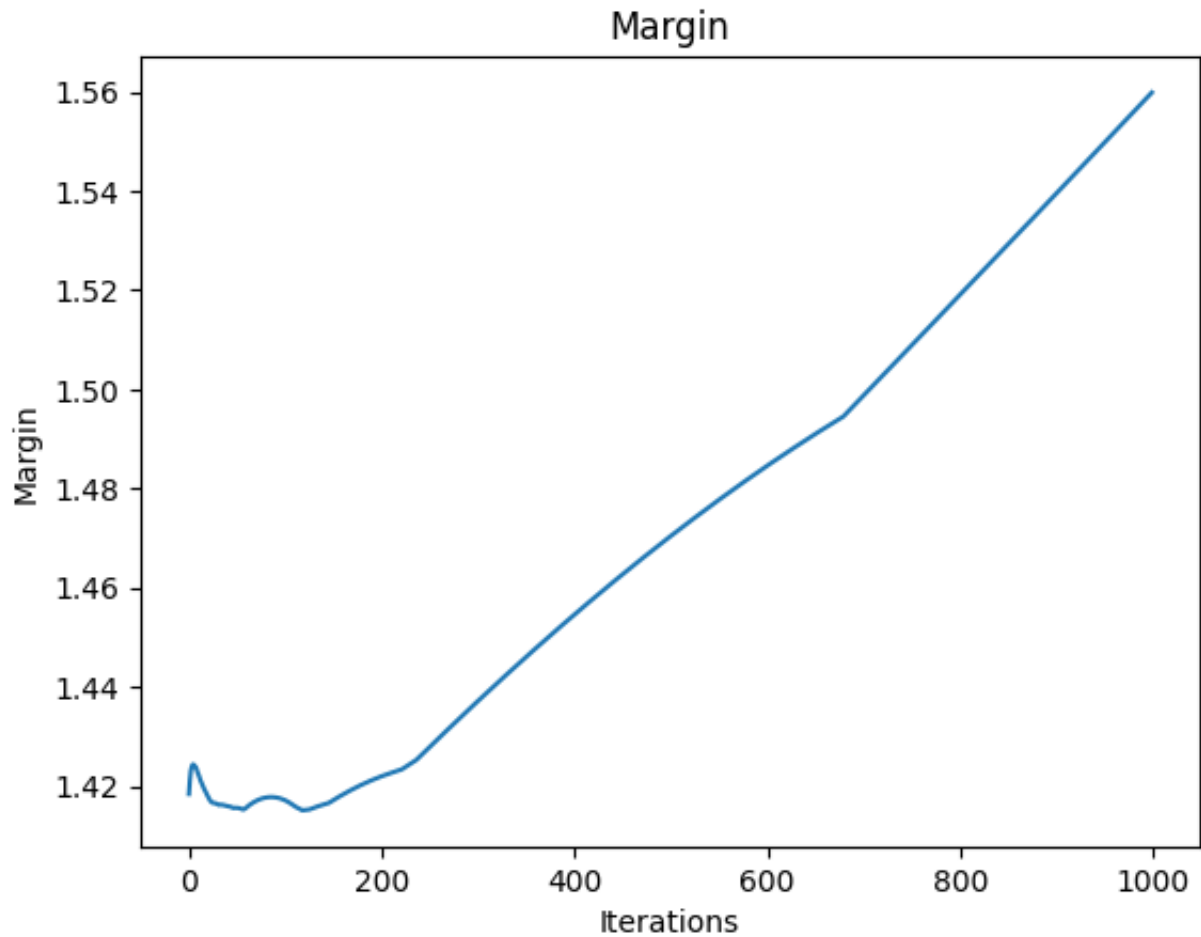
```
In [ ]: plt.plot(error_val)
plt.xlabel('Iterations')
plt.ylabel('Error')
plt.title('Error')
```

```
plt.show()

plt.plot(margin_val)
plt.xlabel('Iterations')
plt.ylabel('Margin')
plt.title('Margin')

plt.show()
```





Part D

Create another scatter plot of your data, but this time color the points by the function $f(x_i) = 1 - y_i \cdot \langle x_i, w \rangle$. The numbers closest to 0 (positive numbers or largest negative numbers) will show you which points were "most important" in determining the classification.

```
In [54]: f_values = [1 - row['y'] * (w[0] * row['x1'] + w[1] * row['x2']) for _, row
plt.figure(figsize=(8, 6))
plt.scatter(data['x1'], data['x2'], c=f_values, label='f(x_i)')
plt.colorbar()
plt.show()
```

