

```
In [48]: # import statements
import numpy as np
from matplotlib import pyplot as plt
from sklearn.datasets import fetch_openml
```

## Question 3

Consider the  $f : \mathbb{R}^2 \rightarrow \mathbb{R}$  with  $f(x_1, x_2) = (x_1 - 2)^2 + (x_2 - 2)^2 - x_1 x_2$  and the following constrained optimization problem:

$$\min_{x_1, x_2} f(x_1, x_2) \quad \text{subject to} \quad 0 \leq x_i \leq 1, \quad i = 1, 2.$$

Write a projected gradient descent algorithm, with constant step size  $\mu = 0.001$  starting at  $(0.5, 0.5)$  for 175 iterations, for the above optimization problem. Plot the function value  $f(x^{(t)})$  against the iteration  $t$ .


```
In [49]: mu = 0.001
x = np.array([0.5, 0.5])
x_values = []
iterations = range(175)

def f(x_1, x_2):
    return (x_1 - 2)**2 + (x_2 - 2)**2 - x_1 * x_2

def df(x_1, x_2):
    return np.array([2 * (x_1 - 2) - x_2, 2 * (x_2 - 2) - x_1])

for i in iterations:
    y = x - mu * df(x[0], x[1])
    x = [min(max(0, y[0]), 1), min(max(0, y[1]), 1)]
    x_values.append(f(x[0], x[1]))

plt.plot(iterations, x_values, marker='.')
plt.xlabel("Iteration")
plt.ylabel(r"Value of $f(x)$")
plt.show()
```



## Question 4

In this coding question, you'll implement a classifier with logistic regression

$$F(w) = \frac{1}{N} \sum_{i=1}^N \log(1 + e^{-\langle w, x_i \rangle y_i}).$$

We will redo the MNIST coding question from HW3.5 but using different versions of gradient descent

$$w^{(t+1)} = w^{(t)} - \mu p^{(t)}.$$

This time, we will differentiate 4's and 9's. That is, instead of classifying images of 0 and 1, you classify the images of 4 and 9.

## Loading MNIST Data

In this section, you will learn to load MNIST data. If you do not have tensorflow available on your jupyter notebook, uncomment the next cell, run it, restart the kernel, and comment the next cell once more.

```
In [50]: # !pip3 install scikit-learn
```

```
In [51]: # this cell will take a minute to run depending on your internet connection
X, y = fetch_openml('mnist_784', version=1, return_X_y=True) # getting data
print('X shape:', X.shape, 'y shape:', y.shape)
```

```
X shape: (70000, 784) y shape: (70000,)
```

```
In [52]: # this cell processes some of the data

# if this returns an error of the form "KeyError: 0", then try running the 1
# X = X.values # this converts X from a pandas dataframe to a numpy array

X = X.values
digits = {j:[] for j in range(10)}
for j in range(len(y)): # takes data assigns it into a dictionary
    digits[int(y[j])].append(X[j].reshape(28,28))
digits = {j:np.stack(digits[j]) for j in range(10)} # stack everything to be
for j in range(10):
    print('Shape of data with label', j, ':', digits[j].shape )
```

```
Shape of data with label 0 : (6903, 28, 28)
Shape of data with label 1 : (7877, 28, 28)
Shape of data with label 2 : (6990, 28, 28)
Shape of data with label 3 : (7141, 28, 28)
Shape of data with label 4 : (6824, 28, 28)
Shape of data with label 5 : (6313, 28, 28)
Shape of data with label 6 : (6876, 28, 28)
Shape of data with label 7 : (7293, 28, 28)
Shape of data with label 8 : (6825, 28, 28)
Shape of data with label 9 : (6958, 28, 28)
```

```
In [53]: # this cell would stack 100 examples from each class together
# this cell also ensures that each pixel is a float between 0 and 1 instead of
data = []

for i in range(10):
    flattened_images = digits[i][:100].reshape(100,-1)
    data.append(flattened_images)

data = np.vstack(data)
data = data.astype('float32') / 255.0
```

## Data PreProcess

```
In [54]: x_4 = digits[4][:500].reshape(500,-1)
x_9 = digits[9][:500].reshape(500,-1)

x_4_test = digits[4][500:1000].reshape(500,-1)
x_9_test = digits[9][500:1000].reshape(500,-1)

x_train = np.vstack((x_4, x_9))
x_train = x_train.astype('float32') / 255.0

x_test = np.vstack((x_4_test, x_9_test))
x_test = x_test.astype('float32') / 255.0

y_train = np.hstack((-1 * np.ones(500), np.ones(500)))
y_test = np.hstack((-1 * np.ones(500), np.ones(500)))
```

## Defining $F(w)$ and $\nabla F(w)$

```
In [55]: def F(w):
    sum = 0
    N = len(x_train)
    for i in range(N):
        sum += np.log(1 + np.exp(-y_train[i] * np.dot(w, x_train[i])))
    return sum / N

def dF(w):
    sum = 0
    N = len(x_train)
    for i in range(N):
        sum += -y_train[i] * np.exp(-y_train[i] * np.dot(w, x_train[i])) * x_train[i]
    return sum / N
```

## $L^2$ Gradient Descent


```
In [56]: mu = 1e-4
T = 100000

w_values = []
iterations = range(T)

w = np.zeros(x_train.shape[1])
for i in iterations:
    w_values.append(F(w))
    w = w - mu * dF(w)

plt.plot(iterations, w_values, marker='.')
plt.xlabel("Iteration")
```

```
plt.ylabel(r"Value of  $F(w)$ ")
plt.show()
```



## Error Rate

```
In [ ]: error = 0

for i in range(1000):
    if np.dot(w, x_test[i]) > 0:
        y_test[i] = 1
    else:
        y_test[i] = -1
    error += (y_test[i] != y_train[i])

print("Error rate:", error / 1000 * 100, "%")
```

Error rate: 4.2 %

## Part 4A

Implement and run  $L^\infty$  gradient descent with step size  $\mu = 10^{-6}$ . Run your algorithm for at least 10000 iterations and initialize with  $w^{(0)} = 0$  (i.e. the zero vector). **Recall:**  $L^\infty$

gradient descent uses steps

$$p^{(t)} = \frac{\text{sign}(\nabla F(w^{(t)}))}{\|\nabla F(w^{(t)})\|_1}.$$


```
In [58]: mu = 1e-6
T = 100000

w_values = []
iterations = range(T)

def p(w):
    f = dF(w)
    return np.sign(f) * np.linalg.norm(f, 1)

w = np.zeros(x_train.shape[1])
for i in iterations:
    w_values.append(F(w))
    w = w - mu * p(w)

plt.plot(iterations, w_values, marker='.')
plt.xlabel("Iteration")
plt.ylabel(r"Value of $F(w)$")
plt.show()
```



## Error Rate

```
In [59]: error = 0

for i in range(1000):
    if np.dot(w, x_test[i]) > 0:
        y_test[i] = 1
    else:
        y_test[i] = -1
    error += (y_test[i] != y_train[i])

print("Error rate:", error / 1000 * 100, "%")
```

Error rate: 5.7 %

## Part 4b

Implement and run  $L^1$  gradient descent (aka. coordinate descent) with step size  $\mu = 10^{-4}$ . Run your algorithm for at least 10000 iterations and initialize with  $w^{(0)} = 0$  (i.e. the zero vector). **Recall:**  $L^1$  gradient descent uses steps

$$p^{(t)} = \frac{\text{sign}(\nabla_{j^*} F(w^{(t)})) \|\nabla F(w^{(t)})\|_{\infty}}{\|\nabla F(w^{(t)})\|_{\infty}} e_{j^*},$$

where  $j^*$  is the location of the largest entry of the gradient, and  $e_j$  is the zero vector with a 1 in the  $j^{\text{th}}$  entry.

```
In [60]: mu = 1e-4
T = 100000


w_values = []
iterations = range(T)
running_sum = np.zeros(x_train.shape[1])

def p(w):
    f = dF(w)
    j = 0
    for i in range(len(f)):
        if abs(f[i]) > abs(f[j]):
            j = i
    running_sum[j] += 1
    return f[j] * np.eye(len(f))[j]

w = np.zeros(x_train.shape[1])
for i in iterations:
    w_values.append(F(w))
    w = w - mu * p(w)

plt.plot(iterations, w_values, marker='.')
plt.xlabel("Iteration")
plt.ylabel(r"Value of $F(w)$")
plt.show()
```





## Error Rate

```
In [61]: error = 0

for i in range(1000):
    if np.dot(w, x_test[i]) > 0:
        y_test[i] = 1
    else:
        y_test[i] = -1
    error += (y_test[i] != y_train[i])

print("Error rate:", error / 1000 * 100, "%")
```

Error rate: 9.700000000000001 %

## Part 4c

Compare these two descent plots of  $F(w)$ , along with the analogous plot for gradient descent from HW3.5. Which performs best, and do you have an argument for why? Do you think the performance would change with different step sizes?

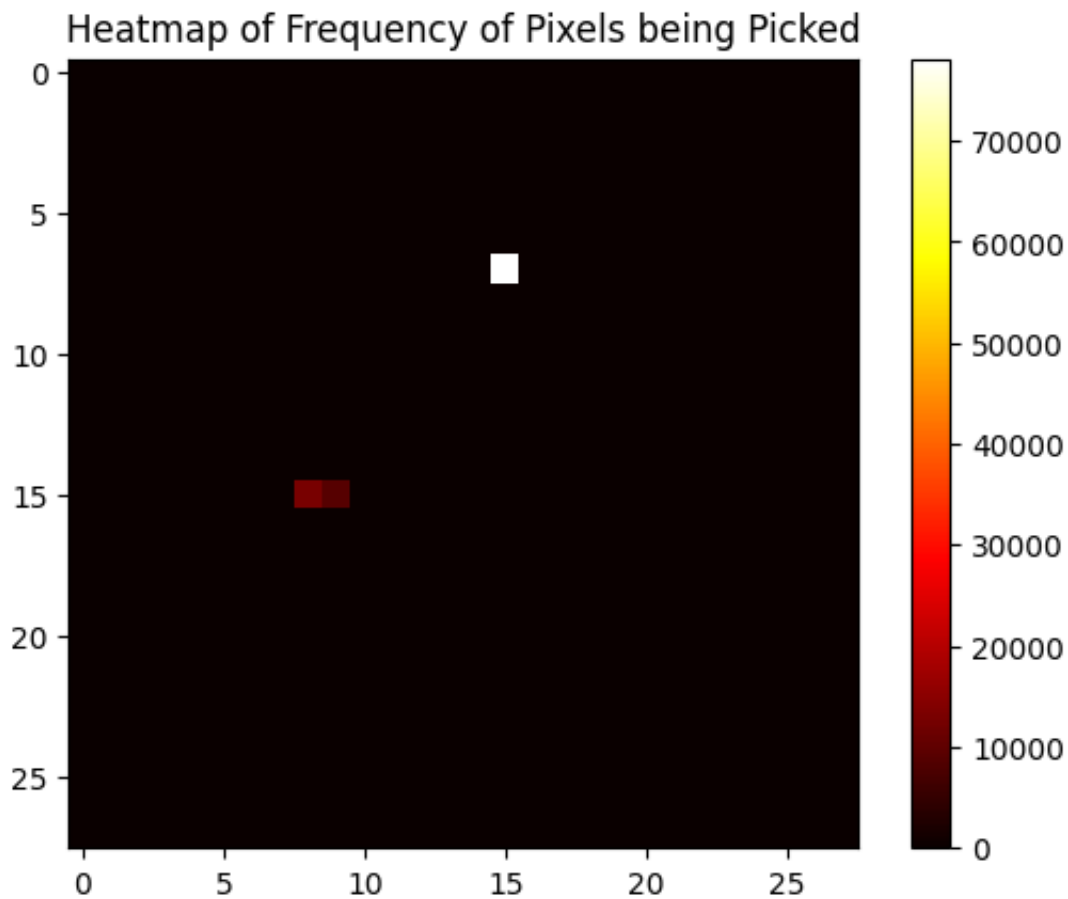
## Answer

Judging from the plots, the  $L^2$  and  $L^\infty$  gradient descent converges the quickest, and the  $L^1$  version converges the slowest, so I would rank their performances in the order of  $L^2 = L^\infty > L^1$ . I believe  $L^1$  gradient descent performs the worst because it only targets one coordinate at a time. On the other hand,  $L^\infty$  and  $L^2$  gradient descent covers all possible descent direction, so they perform better than  $L^1$ . I think step sizes would affect the performances to an extent, but the rank of their performances would not change given the optimal step sizes.

## Part 4d

For the coordinate descent problem, rerun gradient descent but store a running sum of which entry of  $p^{(t)}$  is nonzero at each iteration (not the actual value of the direction vector, just  $e_{j^*}$ ). This will result in a size 784 vector of mostly zeros, and should have integers at various entries whose sum equals the number of iterations. Reshape this vector to be a 28x28 image and display the result. Why do you think these are the pixels that were chosen in the gradient? How can you use this to interpret the algorithm and its results?

```
In [68]: plt.imshow(running_sum.reshape(28, 28), cmap='hot')
plt.colorbar()
plt.title("Heatmap of Frequency of Pixels being Picked")
plt.show()
```



I think these pixels are chosen because they are the most significant features distinguishing 4's and 9's. The most chosen pixels (white and red) are positioned around the area where a 4 would have a shape angle while a 9 would be round. Hence, the  $L^1$  gradient descent, though not being the most performant, can be used to analyze what/where the most influential features are.