

Question 3

In this coding question, you'll implement a classifier with logistic regression

$$F(w) = \frac{1}{N} \sum_{i=1}^N \log(1 + e^{-\langle w, x_i \rangle y_i}).$$

```
In [27]: # import statements
import numpy as np
from matplotlib import pyplot as plt
from sklearn.datasets import fetch_openml
```

Load MNIST Data

```
In [28]: # !pip3 install scikit-learn
# this cell will take a minute to run depending on your internet connection
X, y = fetch_openml('mnist_784', version=1, return_X_y=True) # getting data
print('X shape:', X.shape, 'y shape:', y.shape)
X, y = fetch_openml('mnist_784', version=1, return_X_y=True) # getting data
print('X shape:', X.shape, 'y shape:', y.shape)
# this cell processes some of the data

# if this returns an error of the form "KeyError: 0", then try running the 1
# X = X.values # this converts X from a pandas dataframe to a numpy array

X = X.values
digits = {}
for j in range(10):
    digits[j] = []
for j in range(len(y)): # takes data assigns it into a dictionary
    digits[int(y[j])].append(X[j].reshape(28,28))
digits = {}
for j in range(10):
    digits[j] = np.stack(digits[j])
for j in range(10):
    print('Shape of data with label', j, ':', digits[j].shape )
```

```

X shape: (70000, 784) y shape: (70000,)
X shape: (70000, 784) y shape: (70000,)
Shape of data with label 0 : (6903, 28, 28)
Shape of data with label 1 : (7877, 28, 28)
Shape of data with label 2 : (6990, 28, 28)
Shape of data with label 3 : (7141, 28, 28)
Shape of data with label 4 : (6824, 28, 28)
Shape of data with label 5 : (6313, 28, 28)
Shape of data with label 6 : (6876, 28, 28)
Shape of data with label 7 : (7293, 28, 28)
Shape of data with label 8 : (6825, 28, 28)
Shape of data with label 9 : (6958, 28, 28)

```

Data PreProcess

```

In [29]: x_4 = digits[4][:500].reshape(500,-1)
         x_9 = digits[9][:500].reshape(500,-1)

         x_4_test = digits[4][500:1000].reshape(500,-1)
         x_9_test = digits[9][500:1000].reshape(500,-1)

         x_train = np.vstack((x_4, x_9))
         x_train = x_train.astype('float32') / 255.0

         x_test = np.vstack((x_4_test, x_9_test))
         x_test = x_test.astype('float32') / 255.0

         y_train = np.hstack((-1 * np.ones(500), np.ones(500)))
         y_test = np.hstack((-1 * np.ones(500), np.ones(500)))

```

Define $F(w)$ and $\nabla F(w)$

```

In [30]: def F(w):
         sum = 0
         N = len(x_train)
         for i in range(N):
             sum += np.log(1 + np.exp(-y_train[i] * np.dot(w, x_train[i])))
         return sum / N

         def dF(w):
             sum = 0
             N = len(x_train)
             for i in range(N):
                 sum += -y_train[i] * np.exp(-y_train[i] * np.dot(w, x_train[i])) * x_train[i]
             return sum / N

```

Problem Statement

We will consider the MNIST coding question from HW4. In this question, we run these questions for *differentiating 4's and 9's*. You can reuse the template from the previous homework for loading / formatting MNIST. Implement the following two methods and plot $F(w)$ per iteration for each. You need to submit (i) the code for the algorithms and plots, and (ii) the plots.

Gradient descent with backtracking line search

At each iteration t , initialize the step size $\mu = 10^{-1}$, and use $\gamma = 0.5$ and $\beta = 0.8$ to determine the correct $\mu^{(t)}$. Run your algorithm for at least 10,000 iterations and plot the loss curve $F(w^{(t)})$ as a function of t .

```
In [31]: T = 1000


w_values = []
iterations = range(T)

def gd(w, mu = 1e-1):
    return w - mu * dF(w)

def backtrack(w, gamma=0.5, beta=0.8):
    mu = 1e-1
    while F(gd(w, mu)) > F(w) - gamma * mu * np.linalg.norm(dF(w))**2:
        mu *= beta
    return mu

w = np.zeros(x_train.shape[1])
for i in iterations:
    w_values.append(F(w))
    mu = backtrack(w)
    w = gd(w, mu)

plt.plot(iterations, w_values)
plt.xlabel("Iteration")
plt.ylabel(r"Value of $F(w)$")
plt.show()
```



Error Rate

```
In [32]: error = 0

for i in range(1000):
    if np.dot(w, x_test[i]) > 0:
        y_test[i] = 1
    else:
        y_test[i] = -1
    error += (y_test[i] != y_train[i])

print("Error rate:", error / 1000 * 100, "%")
```

Error rate: 4.3 %

Gradient Descent with Nesterov acceleration.

You can experiment with the parameters until you find something you like.

```
In [35]: T = 1000

w_values = []
```

```


iterations = range(T)

def nesterov(w, w_old, mu = 1e-3, beta = 0.9):
    v = w + beta * (w - w_old)
    return v - mu * dF(v)

w = np.zeros(x_train.shape[1])
w_old = w
for i in iterations:
    w_values.append(F(w))
    w_temp = w
    w = nesterov(w, w_old, mu)
    w_old = w

plt.plot(iterations, w_values)
plt.xlabel("Iteration")
plt.ylabel(r"Value of  $F(w)$ ")
plt.show()

```



Error Rate

```
In [34]: error = 0
```

```
for i in range(1000):  
    if np.dot(w, x_test[i]) > 0:  
        y_test[i] = 1  
    else:  
        y_test[i] = -1  
    error += (y_test[i] != y_train[i])  
  
print("Error rate:", error / 1000 * 100, "%")
```

Error rate: 4.3 %